

IMPACT REPORT

COOL/SCARY AI IMPACT FORUM

San Francisco · November 4, 2025



FROM THE ORGANIZING TEAM

Hello,

Thank you for joining us at the Cool/Scary AI Impact Forum in San Francisco — or for following along from a distance. This report distills the day's conversations so the ideas are easy to share with your team, your board and your community.

San Francisco sits at the geographic center of the AI industry, but the questions we explored were not the ones typically asked in this city. Instead of asking what AI can build, our speakers asked: who gets left behind? Whose bodies bear the cost? Whose voice is missing from the design table? And what would it look like to reclaim agency — as communities, as organizations, as a sector — over how these tools shape our lives?

Five themes emerged. AI's most urgent harms are deeply human. The digital divide has evolved into a question of whether people can make real choices in a society being restructured without their input. The social sector has structural leverage through coalitions and participatory design. Burnout is the operating context for any honest conversation about adoption. And human dialogue is irreplaceable.

How to use this packet: share the session summaries with colleagues, identify one conversation to continue in your organization, and consider what structural change — not just tool adoption — you want to champion this year.

Thanks for showing up, asking hard questions, and building this community with us.

Onward,

Eric Rubin, Ben Childers, and Partners

Cool/Scary AI Impact Forum Organizing Team



THE QUESTIONS WE ASKED

MAIN TAKEAWAY

AI's most consequential impacts are not technical — they are human, structural and political. The social sector is uniquely positioned to respond, but only if organizations move beyond tool adoption to address the systemic conditions that shape who benefits from AI and who is harmed by it.

Start with the people: the burned-out staff, the digitally invisible communities, the women silenced by deepfakes, the data labelers encoding bias at scale. Build coalitions, embed participatory methods, preserve human dialogue, and use the time AI frees up to rest and create — not just to produce more.

This was a room in the industry's home city that declined to have the industry's conversation. The speakers were not debating whether the technology works. They were asking who it works on, and at whose expense — questions that do not get easier the closer you sit to the companies building it.

"Whose bodies bear the cost? Whose voice is missing from the design table?"

THE ORGANIZERS' FRAMING QUESTION



SESSION HIGHLIGHTS

Eric Rubin

CO-FOUNDER, STRATOVATION PARTNERS — CIVIC-SECTOR BENEFITS ENROLLMENT

A case study on See the Stars, using direct-response science normally reserved for fundraising to close the gap between eligible seniors and the services they qualify for: a single intake site reviewing state-specific eligibility rules, multilingual chatbots handling first-layer navigation, predictive models identifying who is at risk of losing coverage before they churn. Eric framed Medicaid churn as a deliberate policy choice, not a systems failure. The tension is trust — which is why partnerships with trusted intermediaries matter.

Jessica Olney

ADVISOR, PEMA — DEEPPAKES AND DEMOCRATIC INTEGRITY

The international human rights lens, drawn from work in Myanmar, Sudan and Yemen. Thousands of state-level AI bills since 2023; California's transparency bills may become seatbelt-level norms. The scarier material: 95 to 99 percent of deepfakes target women, and some tools only work on female bodies. In conflict settings they are used to intimidate women out of public life, with cascading effects on peace processes that depend on whole-of-society participation.

Ben Childers

CO-FOUNDER, STRATOVATION PARTNERS — DIGITALLY INVISIBLE

Argues for a shift from the digital divide, which is input-oriented, to the digitally invisible, which is output-oriented: can people make real choices and fully participate? Grounded in the story of a day laborer who cannot work for the last ten days of every month because his prepaid minutes run out. With billions still lacking reliable access, AI's restructuring risks creating a permanent underclass with no agency over the systems imposed on them.

Aram Fisher

CO-FOUNDER & CEO, CHANGEAGENT — COALITION AI

Big Tech wants users atomized; the social sector thrives on collaboration. ChangeAgent builds shared AI workspaces for coalitions: private instances with models tuned to a coalition's own narrative framework, shared knowledge bases, and natural-language database querying that lets non-technical staff get reporting without bottlenecking data teams. When a model is tailored to your mission, you do not need perfect prompts to get useful output.



SESSION HIGHLIGHTS, CONTINUED

Amanda Bittaker

WIKIMEDIA FOUNDATION — THOUGHTFUL AI INTEGRATION

From the vantage of 200,000 to 300,000 monthly editors curating the world's largest open knowledge base. The scary: AI hallucinates — fabricated ISBNs rose roughly 15 percent after ChatGPT launched — it breaks under adversarial prompting, and it makes disinformation cheaper to produce than to catch. The cool: it cut a multi-minute editorial task to seconds, and 50 non-coders built 20 playable games in a day. Her framework: a human in the loop at every stage, tag all AI-assisted content, reduce toil rather than replace judgment.

Cassie Gruenstein

PLUS IMPACT PARTNERS — AI IN THE AGE OF BURNOUT

Reframes the conversation through chronic occupational burnout: 95 percent of nonprofit leaders report concern about staff burnout, and 60 to 85 percent of knowledge workers identify as burned out. A client produced a board report in 20 minutes instead of 90. The critical move is not to backfill the saved time with more work. Both Keynes in 1930 and Gates today promised technology would reduce work hours; neither delivered. Her prescription is somatic and structural — close stress cycles, change policy, use the sector's labor-advocacy legacy.

Emlyn Metz

LAWRENCE HALL OF SCIENCE, UC BERKELEY — PRESERVING CRITICAL THINKING

The cognitive science of why LLMs are dangerous for epistemology. Their motives are opaque, they pretend to care — driving some users to prefer AI companionship to human connection — and they do not learn from a conversation to carry elsewhere, undermining the model where students co-construct understanding. They exacerbate every vector misinformation exploits: confirmation bias, content overload, polarization, overconfidence without uncertainty signaling. Above all: keep talking to other humans.

Annie Brown

CEO, RELIABLE; AI BIAS RESEARCHER, UC SAN DIEGO — BIAS IN THE LABELING PIPELINE

The most technically grounded talk of the day: bias enters at data annotation, before models are trained. When underpaid overseas workers label images following instructions from data scientists, both parties' cultural biases fuse into the training data. Dark skin tones are represented at just 18 percent at best. One classifier rated a woman doing yoga as racy with 98 percent confidence while a man scored 3. Her solution — participatory labeling — improved accuracy from 67 to 85 percent while reducing bias. The business case and the equity case are the same.



CROSS-CUTTING THEMES

The harms are human, not hypothetical

The most urgent risks are not speculative existential scenarios but concrete, present-tense damage: deepfakes silencing women in conflict zones, classifiers flagging Black women's bodies as explicit, day laborers locked out of work by connectivity gaps, Medicaid recipients churned off coverage by deliberate complexity. The sector's role is not to imagine catastrophe but to respond to what is already happening.

Participatory methods are a structural advantage

Participatory data labeling, coalition workspaces, human-in-the-loop editorial and the community game jam all demonstrated one principle: involving the people affected by AI systems in their design improves both accuracy and trust. This is not a nice-to-have. It is an advantage the social sector is uniquely positioned to hold.

LLMs cannot substitute for human conversation, critical thinking, or the co-construction of understanding. AI collapses pluralism, simulates objectivity, and rewards convenience over reflection. The antidote is not less technology but more deliberate human engagement — which is exactly what events like this exist to provide.

From access to agency

A shift from input metrics — do people have the tools? — to output metrics: can people exercise real choices? The digitally invisible framework, the coalition model prioritizing opt-in autonomy, and the insistence on using AI-freed time for rest all point the same way. The goal is not adoption but self-determination.

Burnout is the operating context

Any AI strategy that ignores the exhaustion of the people who would implement it is incomplete. Every generation's technological promise to reduce work has been absorbed by productivity culture. The opportunity is to use AI to change how we work — reducing toil, creating space for creativity and rest — not to scale existing patterns of overwork.



TENSIONS AND WHAT YOU SAID

Convenience vs. critical thinking

AI makes it easier to produce content and harder to evaluate it. The friction of human writing and dialogue is a feature, not a bug.

Trust vs. speed

Benefits enrollment, coalition data sharing and AI-assisted editorial all require user trust — slow to build, easy to destroy.

Seatbelt vs. patchwork

State-level experimentation is generating useful regulatory models, but the absence of federal frameworks risks compliance chaos.

Atomized users vs. networked power

Big Tech's business model profits from isolated individuals. The social sector's coalition DNA is a direct counter — but requires intentional infrastructure to sustain.

COMMUNITY FEEDBACK

Likelihood to attend again: 4.3 out of 5, with two-thirds giving the maximum score. Attendees valued the scale — small enough to interact with everyone — and the peer connections as much as the stage content. The Coalition AI, Age of Burnout and epistemology sessions were each singled out.

What you asked for: more actionable and technical content alongside the structural framing — what a good AI policy looks like for a nonprofit, which tool for which case. Also structured small-group breakouts, recordings, earlier planning, and a shared link hub for speaker follow-up. A third said yes to the working group and another half asked for more information.



WITH THANKS

THANK YOU TO OUR SPONSORS



CORNERSHOP CREATIVE



ADVOCACY AI



BEEKEEPER GROUP



STRATOVATION PARTNERS

COOL/SCARY AI

A national series of community-centered
conversations about AI's real impact,
produced by Stratovation Partners.

coolslashscary.ai
ben@coolslashscary.ai

